

Out of Distribution
Labs

Hierarchical abstractions as a basis for ontology graphs

Theoretical foundations for software-engineering micro-agents

WHITE PAPER / RESEARCH 01 / VERSION 1.0

Out of Distribution Labs

October 7, 2026

Thesis. A shared graph becomes reasoning infrastructure when abstraction is a declared preservation relationship, evidence remains recoverable, and distributed contributions obey explicit admission rules.

Status. Theoretical synthesis and proposed architecture, supported by reproducible finite checks and a routing simulation. No production deployment, hundred-agent language-model experiment, or new general mathematical theorem is claimed.

<https://outofdistributionlabs.github.io/>

Research process, literature map, evaluation and project history

Abstract

Hundreds of specialized software-engineering agents need a common way to refer to code, requirements, contracts, behavior and evidence. Messages and embedding neighborhoods do not by themselves supply that common semantics. This paper examines hierarchical abstraction as a foundation for ontology graphs, distinguishing semantic abstraction, ontological classification, retrieval organization and execution coordination. Drawing on abstract interpretation, concept lattices, behavioral subtyping, ontology modularity and distributed-state theory, we define contextual abstraction domains and task-relative sufficiency. Elementary propositions establish composition conditions, distinctions that a sufficient abstraction must retain, and a restricted condition for safe parallel vocabulary extension. Counterexamples show why lossy summaries cannot justify universal software claims and why convergent record merging can create an inconsistent ontology. We propose a typed, polyhierarchical graph with revision-scoped claims, abstraction certificates and a controlled admission projection. A reproducible simulation at 100, 300 and 1,000 logical workers illustrates structural routing consequences: a complete multi-view hierarchy matches a complete flat index, missing cross-module membership reduces recall, and coverage-aware fallback increases candidate cost. The framework treats hierarchy as a semantic interface with obligations, not a guarantee that more agents produce better software.

Reader’s guide and claim status

This paper prioritizes theoretical foundations for software engineering. Definitions, propositions and counterexamples are primary; the architecture and simulation are supporting consequences. Propositions follow from explicit assumptions and are not claimed as mathematically novel. The architecture is a proposal. Prior empirical findings retain their original scopes. All measured values concern our synthetic model, not language-model execution. The source register records access depth and versions; the project page links reproducible artifacts.

Contents

1	The problem: shared meaning at agent scale	4
2	Research method and intellectual foundations	4
2.1	Structured narrative review	4
2.2	Abstraction and software specification	4
2.3	Ontologies, evidence and distributed state	5
2.4	Retrieval and agent systems	5
2.5	The synthesis gap	5
3	Four structures, four meanings	6
4	Contextual abstraction domains	6

4.1	Concrete worlds and precision	6
4.2	Composition has explicit premises	7
4.3	Universal properties and spurious counterexamples	7
5	Task-relative sufficiency and polyhierarchy	7
5.1	A cache counterexample	8
5.2	Concept lattices do not certify incidence	8
6	A typed graph and its accepted interpretation	8
6.1	Candidate records versus logical content	8
6.2	Context, identity and multiple parents	9
6.3	An abstraction certificate	9
6.4	Four independent checks	9
7	Safe evolution and distributed admission	10
7.1	Convergence can preserve a contradiction	10
7.2	Conservative extension and parallel vocabulary	10
8	Execution contracts and software workflows	11
8.1	A worked admission trace	11
9	Scaling and out-of-distribution refinement	12
9.1	Arithmetic bounds are not performance laws	12
9.2	OOD as a failure of an interface assumption	12
10	Synthetic illustration: routing and missing membership	13
10.1	Design and measurement	13
10.2	Results and interpretation	13
11	Implementation and empirical research agenda	14
12	Limitations and conclusion	15
A	A minimal certificate interface	16
B	Finite checks and reproduction	16

C Reusable research process

16

1 The problem: shared meaning at agent scale

A repository contains multiple structures: files and packages, type hierarchies, call graphs, data flow, runtime dependencies, requirements and change histories. None is a complete description of the others. A function can belong to one package, implement a security-sensitive operation, and participate in several deployment workflows. A module summary may erase precisely the behavior another worker needs to verify.

Distributed specialization increases the pressure on shared meaning. A parser identifies symbols, a contract worker describes preconditions, a security worker follows effects, a test worker constructs witnesses, a patch worker changes code, and a reviewer checks evidence. Their outputs must refer to compatible revisions and vocabulary while retaining different authority and scope. Sharing a vector store does not prevent evidence about one environment from being used as a fact about another.

We use *micro-agent* to mean a bounded worker with a declared role, tools, readable vocabulary and admissible effects. It may be a language-model process, deterministic analyzer, or human-assisted worker. Hundreds of registered workers need not mean hundreds of concurrent model calls or separate models. A large registry can supply a small task-specific cohort.

The central question is semantic: *what must an abstraction preserve so another worker can reason through it?* Efficiency is a separate question. We propose that an ontology graph should expose preservation, grounding and context, with retrieval and scheduling attached through separate interfaces. The graph does not replace compiler semantics, testing or execution control.

2 Research method and intellectual foundations

2.1 Structured narrative review

We reviewed 22 primary-source records with a cutoff of October 7, 2026. Search families covered abstraction, concept lattices, ontology classification and modularity, software specification, graph retrieval, multi-agent coordination and distributed convergence. Inclusion required a directly relevant mechanism or formal distinction. We preferred original papers, author/publisher records, conference proceedings and normative W3C specifications; exact queries and access notes are published.

This is not an exhaustive systematic review. Several foundational records were inspected at author-summary, abstract or bibliographic level rather than read in full: abstract interpretation's author summary, the concept-analysis book, Parnas, Contract Net, the blackboard overview and SWE-bench. Hoare's scanned PDF yielded no usable local text, so its cited role is limited to the inspected primary-source passage. Full-text readings included behavioral subtyping, OntoClean's revised chapter, ontology modularity, the CRDT report, GraphRAG, HippoRAG, AutoGen, multi-agent failure analysis, agent scaling and SWE-agent. Adjacent source-code knowledge graphs and MDP abstraction were screened but are not relied on for propositions.

2.2 Abstraction and software specification

Abstract interpretation connects simplified representations with concrete program semantics, allowing imprecision while preserving soundness [2]. Formal concept analysis orders closed extent-intent pairs

induced by an object–attribute incidence relation [4]. The first concerns interpretation and preservation; the second concerns structure induced by declared attributes. Neither makes an arbitrary generated summary a proof.

Parnas distinguishes criteria for decomposing a system rather than treating one convenient decomposition as inevitable [13]. Behavioral subtyping constrains substitution through specifications, invariants and relevant histories [11]. Hoare-style pre/postcondition reasoning makes program obligations explicit [8]. These foundations suggest question-relative representations and explicit contracts, not a universal tree of code concepts.

2.3 Ontologies, evidence and distributed state

OntoClean uses metaproperties including identity, rigidity and unity to analyze taxonomic choices [6]. Ontology modularity formalizes safe reuse through conservative extension and supplies locality-based sufficient mechanisms; general problems can be undecidable in expressive settings [5]. OWL 2 profiles define different reasoning fragments, and Direct Semantics defines logical interpretation [17, 18].

SKOS supplies knowledge-organization relations, SHACL validates selected graph constraints, and PROV-O represents derivation and attribution [16, 19, 20]. These are different obligations. CRDT theory supplies convergence conditions for suitable replicated state [14]; convergence does not certify the logical consistency of the accumulated statements.

2.4 Retrieval and agent systems

GraphRAG uses hierarchical community summaries for global query-focused summarization [3]. HippoRAG integrates knowledge graphs and Personalized PageRank for evidence retrieval [7]. Topical membership and relevance scores are not certified class inclusion or truth.

Contract Net separates task announcement, bidding and award, while blackboard systems coordinate specialized knowledge sources through shared state [12, 15]. AutoGen supplies programmable conversation patterns [21]. Multi-agent failure analysis distinguishes specification/design, alignment, verification and termination failures [1]; scaling comparisons show dependence on task and architecture, not uniform gains from adding workers [10]. None establishes hundred-agent behavior for our proposal.

SWE-agent studies interfaces for repository navigation, editing and testing; SWE-bench supplies executable repository-level issue tasks [9, 22]. They inform a future evaluation plan, rather than validate the ontology architecture here.

2.5 The synthesis gap

These mechanisms answer different questions: how to simplify semantics, organize concepts, retrieve evidence, allocate tasks and transport records. Our contribution is a proposed interface between them, not a claim that they have never been combined. An abstraction should expose preservation assumptions, a retrieved summary should expose evidence, and an accepted assertion should expose admission scope. Whether this combination improves software outcomes remains untested.

3 Four structures, four meanings

Structure	Relationship represented	What it does not certify
Semantic abstraction	Simpler representation with preserved properties	Correctness of unspecified compression
Ontology	Classes, instances, predicates and axioms	Behavioral substitution from names
Retrieval organization	Communities, summaries and relevance	Logical implication or truth
Coordination	Dependencies, capabilities and effects	Correctness of worker conclusions

Table 1: One graph can store all four structures, but its edge types and inference rules must remain distinct.

A payment function is *part of* a package, *calls* a retry utility, and may have a contract that *requires* an idempotency key. A tool report *supports* a behavior claim; a module summary *describes* related records; a task is *assigned to* a patch worker. Replacing these relations by a generic parent edge destroys the conditions under which an inference is allowed.

Hierarchy therefore means a family of ordered views. A taxonomic view may be a DAG after equivalent classes are collapsed. Containment can be a tree or DAG. Call and dependency graphs can be cyclic. Coordination can have feedback even if one scheduled task run is acyclic. The representation must declare these choices rather than import assumptions from a visualization.

4 Contextual abstraction domains

4.1 Concrete worlds and precision

Fix $c = (r, e, v)$: repository revision, environment and specification version. Let X_c denote admissible concrete software states or traces. A concrete knowledge state is $S \in C_c = \mathcal{P}(X_c)$ ordered by inclusion. More possible worlds means less precision. This is a semantic model, not a claim that all program executions are enumerable.

Let $(A_c, \sqsubseteq)$ be an abstract ordered domain with monotone abstraction and concretization maps $\alpha : C_c \rightarrow A_c$, $\gamma : A_c \rightarrow C_c$. In the standard adjunction setting,

$$\alpha(S) \sqsubseteq a \iff S \subseteq \gamma(a). \quad (1)$$

Consequently $S \subseteq \gamma(\alpha(S))$. The abstraction may include spurious possibilities but must not exclude an actual one. Completeness and computability are separate. This construction derives from established abstraction theory [2]; applying it to a graph requires real interpretation maps.

For integer possibilities $S = \{2, 4\}$, an interval abstraction gives $[2, 4]$. It also permits 3, but does not say 3 was observed. Likewise, “may acquire a lock” does not mean “always acquires a lock.” Workers need the modality of a representation, not merely its label.

For a concrete transformer $F : C_c \rightarrow C_c$, a sound abstract transformer $F^\# : A_c \rightarrow A_c$ satisfies

$$F(\gamma(a)) \subseteq \gamma(F^\#(a)) \quad \text{for all } a. \quad (2)$$

A concise LLM explanation has no such guarantee by virtue of its fluency. It can be a candidate view or a retrieval aid until its obligations are checked.

4.2 Composition has explicit premises

Proposition 1 (Composition of registered abstractions). *If ordered domains A_0, A_1, A_2 have Galois connections $(\alpha_{01}, \gamma_{10})$ and $(\alpha_{12}, \gamma_{21})$, their composed maps also form a Galois connection.*

Proof. For $x \in A_0, z \in A_2$,

$$\begin{aligned} \alpha_{12}(\alpha_{01}(x)) \sqsubseteq_2 z &\iff \alpha_{01}(x) \sqsubseteq_1 \gamma_{21}(z) \\ &\iff x \sqsubseteq_0 \gamma_{10}(\gamma_{21}(z)). \end{aligned}$$

These are precisely the adjunction conditions for the composite. □

This elementary result licenses compatible, registered paths. It does not prove $\text{code} \rightarrow \text{module} \rightarrow \text{service} \rightarrow \text{business goal}$ is a sound chain. A topical summary, a class relation and a behavior projection need not share domains, context or preservation properties.

4.3 Universal properties and spurious counterexamples

Let $\llbracket \varphi \rrbracket_c$ be the concrete worlds satisfying φ . If $S \subseteq \gamma(a)$ and $\gamma(a) \subseteq \llbracket \varphi \rrbracket_c$, then $S \subseteq \llbracket \varphi \rrbracket_c$ by transitivity. A universal property proved for all concretizations of a sound abstraction applies to the actual represented possibilities.

The reverse fails: a coarse abstraction can introduce a spurious violating world even when all actual worlds satisfy the property. A violation of the coarse model calls for a concrete witness or refinement, not an automatic bug report. Conversely, favorable observations do not prove a universal premise without a coverage argument. A passing finite test suite and a sound analysis carry different guarantees.

5 Task-relative sufficiency and polyhierarchy

Definition 1 (Question equivalence and sufficiency). *For Boolean predicates Q on X_c , define $x \sim_Q y$ when $q(x) = q(y)$ for all $q \in Q$. A map $h : X_c \rightarrow Z$ is Q -sufficient when every q factors through h : $q = \bar{q} \circ h$ for some $\bar{q} : Z \rightarrow \{0, 1\}$.*

Proposition 2 (Necessary separation). *If some $q \in Q$ distinguishes x from y , every Q -sufficient map distinguishes them.*

Proof. If $h(x) = h(y)$, then $\bar{q}(h(x)) = \bar{q}(h(y))$, contradicting $q(x) \neq q(y)$. □

For finite X_c , the quotient $X_c / \sim_Q$ is a sufficient partition; every sufficient partition refines it. This gives a lower bound on the distinctions the question family demands. It neither makes equivalence decidable for arbitrary program behavior nor provides an automatic way to infer all relevant questions from a request.

5.1 A cache counterexample

Two cache implementations have identical sequential outputs. One invalidates atomically; the other permits a concurrent reader to observe stale data after invalidation. A summary retaining only sequential input/output behavior maps them to the same representation. It is insufficient for “can a stale read occur after invalidation?” It may remain sufficient for a recorded sequential test question.

API, effect, security and deployment views should therefore coexist. An interface-discovery view can be inadequate for concurrency verification. Refinement can preserve one necessary distinction without expanding every summary into the entire codebase. The graph should retain which question family a view serves and the witness that caused a split.

5.2 Concept lattices do not certify incidence

Formal concept analysis begins with (O, M, I) : objects, attributes and incidence. Derivation operators yield closed extent–intent pairs ordered by extent inclusion [4]. This supports overlapping classification without requiring a tree. It does not establish the truth of I .

In software, “passed retry test” and “may access credentials” have different modalities. Measured attributes need revision and environment scope; possible effects must not be turned into recorded runtime occurrences. A lattice built from mistaken incidence is orderly but mistaken. We use it as organization of declared information, with evidence admission treated separately.

6 A typed graph and its accepted interpretation

6.1 Candidate records versus logical content

Let $G_r = (V, E, \tau, \kappa, P)$ be a snapshot: nodes and edges, sorts τ , contexts κ , and provenance links P . Let T_r be a separate set of admitted axioms. Adding a candidate claim record to V does not automatically add its asserted formula to T_r .

Useful sorts include Artifact, Concept, Requirement, Contract, Observation, Claim, Task, Capability and Context. An ontology class is not a source-code object with the same printed name. A test run is an observation, not a proof by default. Predicate signatures expose these distinctions operationally.

Predicate	Intended relation	Inference boundary
subClassOf	Logical class inclusion	Compatible interpretation and context
partOf	Composition or containment	Does not imply subclass membership
calls / dependsOn	Dependency	Does not transfer all properties
summarizes	Grounded view	No entailment without preservation conditions
supports / derivedFrom	Evidence and provenance	No automatic truth or independence
assignedTo	Allocation	No verification guarantee

Table 2: Illustrative relation contracts; an implementation must provide actual schemas and semantics.

6.2 Context, identity and multiple parents

A function can be part of a payment module, meet an idempotency contract, and participate in a sensitive-effect view. These are different paths. For this application we propose acyclic taxonomic views after declared equivalences are collapsed. This is an application invariant, not a universal constraint on OWL or SKOS. Other graph relations can remain cyclic.

Use revision-qualified identifiers, source hashes and explicit rename/move lineage. Matching paths do not prove identity: a symbol can be deleted and recreated at the same location. A move can preserve identity while invalidating location-based evidence. Identity mappings themselves need provenance and review.

OntoClean’s stable-kind versus contextual-role distinction informs this design [6]. “Assigned reviewer” is a role, not necessarily a permanent agent superclass. “Currently failing test” refers to an observation under a revision and environment, not an intrinsic type of a file.

6.3 An abstraction certificate

A promoted abstraction relationship should expose: source/target sorts; concrete context; interpretation maps or an explicitly weaker relation; preserved question family or property fragment; evidence witnesses and proof obligations; validation method; source hashes; and invalidation conditions. A missing obligation does not forbid storing the view: it changes its status to hypothesis or retrieval aid.

This certificate is a proposed interface, not an assertion that all obligations are automatically decidable. Restricted analyses can discharge some formally; informal requirements may require review. “Proof-backed,” “validated on selected tests” and “unverified summary” should remain distinguishable in every downstream use.

6.4 Four independent checks

Class inclusion $A \subseteq B$ transfers universal properties of B to A under one interpretation. It does not transfer an arbitrary existential witness in B into A , and it does not prove implementation substitutability. Behavioral substitution concerns preconditions, postconditions, invariants and relevant history constraints [11].

OWL open-world entailment and SHACL validation must remain separate [18, 20]. A missing required field can fail a selected shape without entailing that a real-world relation is false. SKOS broader is not itself declared transitive, and broaderTransitive is distinct [16]; a knowledge-organization edge must not silently become a subclass axiom.

PROV records origin and derivation [19], not truth, calibrated confidence or source independence. Ten summaries derived from one test remain one underlying observation. Authorization is a fourth concern: semantic relevance to a file does not grant a worker permission to edit it.

7 Safe evolution and distributed admission

7.1 Convergence can preserve a contradiction

CRDT convergence applies under declared algebraic and delivery assumptions [14]. Consider T_0 declaring A and B disjoint, with proposals $\Delta_1 = \{A(a)\}$ and $\Delta_2 = \{B(a)\}$. Each extension of T_0 has a model, but $T_0 \cup \Delta_1 \cup \Delta_2$ does not. Claim records can converge perfectly while their admitted interpretation becomes inconsistent. Record convergence and semantic admission are independent obligations.

We propose an append-only candidate log and a controlled accepted projection. Identity merges, competing contracts and schema edits are non-monotone choices requiring checks or serialized governance. Retractions should preserve supersession/tombstone history. A deterministic admission pipeline can check authorization, signatures, shapes, revision dependencies, selected logical consistency conditions and evidence policy. It should record which criteria passed rather than flatten them into “verified.”

7.2 Conservative extension and parallel vocabulary

For signature Σ , a deductive conservative extension preserves old Σ entailments:

$$\forall \varphi \text{ over } \Sigma, \quad T' \models \varphi \iff T \models \varphi. \quad (3)$$

A model-conservative extension is stronger: every model of T can be expanded, on the same domain and with old symbols unchanged, to a model of T' . These notions underlie ontology modularity [5].

Proposition 3 (Independent private vocabulary extension). *Let T use signature Σ . For $i = 1, 2$, let $T_i = T \cup \Delta_i$ be model-conservative over T using private signatures Σ_i disjoint from Σ and from each other. Suppose Δ_i mentions only $\Sigma \cup \Sigma_i$. Then $T \cup \Delta_1 \cup \Delta_2$ is model-conservative over T .*

Proof. Take a model M of T . Each premise supplies an expansion M_i on the same domain, preserving M on Σ . Interpret the private symbols in Σ_1 as in M_1 and those in Σ_2 as in M_2 . Disjointness makes this combined interpretation well-defined. Every axiom of Δ_i sees exactly its interpretation from M_i , so both sets hold. The combination expands the arbitrary model M of T . $\square$

This elementary sufficient condition explains a limited form of parallel semantic work. It does not license overlapping contract edits, shared identity changes or arbitrary claim additions. The disjointness counterexample fails the conservativity premises. Ordinary code ownership does not establish them either.

General conservativity checking can be undecidable in expressive ontology settings. Locality extraction provides useful sufficient mechanisms, but a nearby graph neighborhood is not automatically a logical module [5]. Restricted profiles, extracted modules and finite regression query suites are practical alternatives. Passing a finite suite certifies that suite, not full conservativity. Intentional corrections may legitimately change old entailments; those changes should be explicit and communicated.

8 Execution contracts and software workflows

Represent a worker interface as

$$a = (\rho, \Sigma_R, W, \mathcal{U}, Pre, Out, Verify, B), \quad (4)$$

where the entries are role, readable signature, write scope, tools, preconditions, output shapes, verification policy and budget. A worker reads an immutable snapshot, selects a relevant view, expands to grounded artifacts, performs a bounded action and proposes a delta with evidence.

Read dependencies matter as much as edited files. Two workers can edit different files while one depends on a contract the other changes. Revision checks should cover the relevant dependency closure or mark the result stale. A graph records these dependencies; scheduling and execution isolation enforce them.




Figure 1: Proposed interfaces. Candidate storage is not semantic acceptance. Dashed feedback exposes accepted snapshots to retrieval.

For safe concurrent payment retries, a requirements worker records the intended property and ambiguities. Parser/effect workers ground symbols and possible side effects. A contract worker identifies idempotency obligations. A counterexample worker seeks violating interleavings. A patch worker proposes a revision; verification workers run selected analyses and tests. Admission checks the evidence and policy before producing a new accepted snapshot. Security and integration concerns may cut across package ancestry; the workflow requires a disagreement path, not only a success path.

Hundreds of workers can share vocabulary while only a bounded cohort executes a task. Their role and relevance do not grant identical write authority. Textual consensus is not a proof, and all tools may share the same mistaken assumptions. Failure patterns reported in multi-agent systems motivate explicit verification and termination policies, not an assumption that a graph automatically removes them [1].

8.1 A worked admission trace

Suppose revision r_0 has a function $\text{charge}(\text{key}, \text{amount})$ and a requirement that concurrent retries must not create duplicate ledger entries. A module summary says “the payment service is idempotent.” That sentence is only a candidate claim until its modality and evidence are resolved.

The parser grounds the function and ledger operation. A unit-test worker reports two sequential retries yielding one entry in environment e_0 . Its claim is an observation about those executions. A contract worker proposes the stronger predicate that all admitted concurrent traces contain at most one entry for a key. An effect worker finds a check-then-write path; a counterexample worker constructs two

overlapping executions that both see an absent key. This witness separates traces the original summary identified.

The accepted projection may admit the unit-test observation while refusing the universal contract. It retains the contradiction witness and marks the old summary insufficient for the concurrency question. A patch worker proposes an atomic ledger insertion at r_1 ; verification runs at (r_1, e_0, v_0) , not silently at the old context. Passing a new stress test is still empirical evidence, while a proof about the atomic operation depends on its specified semantics and environment assumptions.

This trace explains what the hierarchy contributes: navigation from requirement to contract to grounded operations, and a location to refine a failed abstraction. It does not decide the atomicity of an external database by itself. If a later environment changes the transaction semantics, the certificate's dependency invalidates even if the source function is unchanged.

9 Scaling and out-of-distribution refinement

9.1 Arithmetic bounds are not performance laws

A complete directed all-to-all topology has $N(N - 1)$ possible edges. A policy notifying at most k workers creates at most k deliveries per update. For balanced disjoint buckets, selecting b of B buckets contains approximately Nb/B workers. These are counting statements with assumptions, not a theorem that ontology use makes total coordination linear.

Skew, overlapping expertise, broad requests and cross-cutting effects alter cohort size. A flat inverted index with the same memberships can select the same workers. Index maintenance, validation, replication, rework and model execution remain separate costs:

$$C_{\text{task}} = C_{\text{route}} + C_{\text{grounding}} + C_{\text{workers}} + C_{\text{validation}} + C_{\text{rework}}. \quad (5)$$

Graph construction/maintenance must additionally be measured or amortized under a stated workload. Agent-scaling evidence motivates testing task–architecture alignment [10], not claiming a universal savings factor.

9.2 OOD as a failure of an interface assumption

Software OOD conditions include unfamiliar APIs, unrepresented effects, missing schema relations, changed dependencies, incompatible environments, stale evidence and counterexamples to a contract. A novel embedding does not calibrate the chance an inference is wrong.

Separate four flags: outside-schema evidence; unsupported preservation step; insufficient evidence coverage; and logical conflict. A flag should identify an affected view/claim and trigger grounding, refinement, abstention or review. It should not automatically induce a global schema change.

A refinement procedure retains the witness, locates the representation that merged distinguishable cases, proposes a local split or relation, reruns affected checks and records supersession. A concurrency witness may require a trace/effect view, not a source-code superclass. Broadening a concept to absorb every exception can make its preserved properties vacuous.

10 Synthetic illustration: routing and missing membership

10.1 Design and measurement

Our Python simulation uses 100, 300 and 1,000 logical workers, 16 modules with four skills each, 20 seeds (0–19), and 300 generated requests per seed/condition. Each worker receives a primary module/skill label and a secondary label in another module with probability 0, .25 or .5. Labels also define eligibility truth.

Baselines are broadcast, complete flat index, primary-only hierarchy and guarded complete multi-view hierarchy. Local queries select one module/skill. Cross-cutting queries select one skill in four modules. Coverage-gap queries select a local label while each label's coverage flag is independently false with probability .1; the guard then broadcasts. Coverage detection is an oracle in this model, not an implemented OOD detector.

We measure candidates, eligible-worker recall and structural index visits. Visits count selected nodes and coverage checks, not runtime latency. Candidates can be interpreted as notification counts only for a policy sending one notification per candidate. No LLM is called, patch generated or token saving measured.

Scenario	Strategy	Candidates	Recall	Visits
local	broadcast	300.00	1.000	0.0
local	flat index	7.07	1.000	1.0
local	primary hierarchy	4.69	0.664	4.0
local	guarded multiview	7.07	1.000	5.0
cross cutting	broadcast	300.00	1.000	0.0
cross cutting	flat index	27.73	1.000	4.0
cross cutting	primary hierarchy	18.80	0.678	11.9
cross cutting	guarded multiview	27.73	1.000	15.9
coverage gap	broadcast	300.00	1.000	0.0
coverage gap	flat index	7.03	1.000	1.0
coverage gap	primary hierarchy	4.67	0.663	4.0
coverage gap	guarded multiview	38.71	1.000	5.0

Table 3: Means at 300 workers and .5 secondary-membership probability, averaging 20 seed-level means. Recall excludes empty eligible sets. Full raw records, retained counts and seed standard deviations are published.

10.2 Results and interpretation

Local requests select 7.07 workers on average under both complete strategies, versus 300 under broadcast. The primary-only hierarchy selects 4.69 but retains only .664 eligible-worker recall. Cross-cutting requests select 27.73 under the complete strategies, versus 18.80 and .678 recall under primary-only membership. For coverage gaps, guarded recall stays 1.000 but candidate count rises to 38.71; the complete flat baseline selects 7.03. These are generated-label results, not measurements of software competence.

The simulation illustrates a distinction visible in the formal model: a smaller cohort may result from missing information rather than successful abstraction. Complete hierarchy and flat indexing produce




Figure 2: Candidate fractions at 300 workers. Complete flat-index and multi-view curves coincide without fallback. Broadcast is a weak routing baseline, not evidence that hierarchy outperforms an equivalent index.




Figure 3: Local recall. Loss results from omitted secondary membership in the primary-only ablation. Complete strategies coincide at 100%; this is not an intrinsic failure of every hierarchy.

the same candidate sets without fallback. Hierarchy's potential semantic advantage is not established by these counts.

The raw data contains 2,160 seed/strategy records. Uncertainty is sample standard deviation of seed-level means, not a confidence interval. Empty eligible sets are excluded from recall and retained counts are published; conditioning affects small- N averages. Input and truth use the same generator, so complete-index recall is guaranteed by construction. Maintenance, stale state, parser errors, actual messages, contention and patch quality are outside the model. Perfect coverage detection is deliberately favorable and requires separate empirical validation.

11 Implementation and empirical research agenda

Start with a bounded vocabulary of artifacts, contracts, claims and contexts, deterministic symbol extraction, separate candidate/accepted projections, immutable snapshots and structured delta proposals. Use a declared reasoning fragment rather than unrestricted generated axioms. OWL profile properties

constrain logical work but do not guarantee low end-to-end latency [17].

A future study should compare single-agent execution, unstructured multi-agent memory, a complete flat typed index, hierarchy without certificates, and hierarchy with certificates/admission. Match models, tools, repository snapshots, evidence and task budgets. Distinguish registry size from concurrently active cohort size. A 1,000-worker registry with five active workers is not a concurrent 1,000-agent experiment.

Tasks should include bug localization, change impact, security analysis and patch generation. SWE-bench-style executable tasks offer one starting point [9]; supplement test-defined success with hidden tests, static checks, domain obligations and review. Keep development tasks separate from held-out evaluation; stress unseen repositories, dependency changes and stale observations.

Measure verified outcomes, unsupported-claim rate, evidence recall, stale-read rejection, conflict detection, provenance recovery and cost per verified result. Separate wall-clock/model-call costs from graph operations. Include graph construction and update overhead. A hierarchy that improves explanation but not patch success may still have value, but it should be evaluated against that stated purpose.

Theoretical questions remain: how to infer a useful question family from requirements; which certificates admit tractable checks; how to align multiple views after code changes; and how to detect coverage gaps without equating novelty with error. An elegant ontology cannot repair incorrect concrete semantics; a useful grounded analyzer need not require a rich ontology.

12 Limitations and conclusion

This paper supplies a synthesis, definitions, elementary proofs, counterexamples and a proposed interface. It does not establish mathematical novelty, a complete software ontology, an implemented semantic reasoner or empirical improvement for hundreds of LLM agents. The review is source-traceable but not exhaustive; limited access to some foundational sources is disclosed. Simulation results concern generated labels, not a deployment.

The central principle is to make abstraction a declared relationship with preservation obligations. Distinct views answer distinct questions; typed relations prevent association from masquerading as entailment; revision scope prevents evidence from crossing contexts silently; and admission separates convergent storage from accepted theory. These foundations make an ontology graph a plausible semantic interface for many specialized workers. Its practical value must be tested through verified software outcomes with strong simple baselines and explicit failure cases.

A A minimal certificate interface

Field	Meaning
Claim identity	Immutable record distinct from its asserted formula
Context	Revision, environment, specification and viewpoint
Typed assertion	Predicate with declared source/target signature
Evidence status	Hypothesis, observation, checked obligation or proof-backed claim
Evidence lineage	Source hashes, tool run, conditions and provenance
Preserved questions	Property fragment or question family covered
Dependencies	Read symbols/certificates invalidated on change
Admission	Accepted, disputed, rejected or superseded, with reason

Table 4: Proposed interface. Confidence scores require their own calibration method and evaluation.

A test report states revision, environment and tested conditions. A universal contract states proof or abstraction assumptions. A patch worker must not upgrade the former into the latter to fill a schema. This discipline can be implemented before a complex ontology is adopted.

B Finite checks and reproduction

The general proofs appear in the main text. The supporting script enumerates a finite concrete domain:

$$X = \{-2, -1, 0, 1, 2\}.$$

It checks its powerset, intervals and coarser intervals with endpoints $\{-2, 0, 2\}$. It checks 512 concrete–interval adjunction instances, 112 interval–coarse instances and 224 composed instances, plus the summary and merge counterexamples. These are finite sanity checks, not machine proofs for arbitrary programs.

From the project folder:

```
python evaluation/simulate.py
python evaluation/check_formal.py
python evaluation/plot.py
latexmk -pdf white-paper.tex
```

Simulation/checking use Python’s standard library; plots require Matplotlib. The PDF uses pdfLaTeX/BibTeX and ordinary TeX Live packages. Raw CSV, aggregate JSON, plots, notes and source records are published. Bibliographic years reflect original publication/preprint dates with pinned revisions noted in the bibliography.

C Reusable research process

The shared runner tools/research/run.py reads project configuration, renders a static progress page, commits explicit project paths, pushes and sends ntfy. A two-minute heartbeat refreshes status while running; it does not itself perform research or independently certify progress. Completion stops it.

The six stages are protocol, literature, synthesis, evaluation, manuscript review and publication verification. The template requires scope confirmation, source/version tracking, claim-status distinctions, appropriate baselines, reproducible evidence and final artifact checks. It is a reusable process rather than a prescription for the answer.

References

- [1] Mert Cemri, Melissa Z. Pan, Shuyi Yang, Lakshya A. Agrawal, Bhavya Chopra, Rishabh Tiwari, Kurt Keutzer, Aditya Parameswaran, Dan Klein, Kannan Ramchandran, Matei Zaharia, Joseph E. Gonzalez, and Ion Stoica. Why Do Multi-Agent LLM Systems Fail? arXiv:2503.13657v3, 2025. URL <https://arxiv.org/abs/2503.13657v3>. Failure-analysis preprint; accessed October 7, 2026.
- [2] Patrick Cousot and Radhia Cousot. Abstract Interpretation: A Unified Lattice Model for Static Analysis of Programs by Construction or Approximation of Fixpoints. In *Proceedings of the Fourth ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, pages 238–252, 1977. URL <https://cs.nyu.edu/~pcousot/COUSOTpapers/POPL77.shtml>. Foundational theory; accessed October 7, 2026.
- [3] Darren Edge, Ha Trinh, Newman Cheng, Joshua Bradley, Alex Chao, Apurva Mody, Steven Truitt, Dasha Metropolitan, Robert Osazuwa Ness, and Jonathan Larson. From Local to Global: A Graph RAG Approach to Query-Focused Summarization. arXiv:2404.16130v2, 2024. URL <https://arxiv.org/abs/2404.16130v2>. Graph retrieval preprint; accessed October 7, 2026.
- [4] Bernhard Ganter and Rudolf Wille. *Formal Concept Analysis: Mathematical Foundations*. Springer, 1999. URL <https://doi.org/10.1007/978-3-642-59830-2>. Foundational theory; accessed October 7, 2026.
- [5] Bernardo Cuenca Grau, Ian Horrocks, Yevgeny Kazakov, and Ulrike Sattler. Modular Reuse of Ontologies: Theory and Practice. *Journal of Artificial Intelligence Research*, 31:273–318, 2008. doi: 10.1613/jair.2444. URL https://www.cs.ox.ac.uk/oucl/work/yevgeny.kazakov/publications/journ/CueHorKazSat08Modularity_JAIR.pdf. Ontology theory; accessed October 7, 2026.
- [6] Nicola Guarino and Christopher A. Welty. An Overview of OntoClean. Handbook on Ontologies, second edition, 2009. URL <https://www.loa-cnr.it/Papers/GuarinoWeltyOntoCleanv3.pdf>. Ontology methodology; accessed October 7, 2026.
- [7] Bernal Jiménez Gutiérrez, Yiheng Shu, Yu Gu, Michihiro Yasunaga, and Yu Su. HippoRAG: Neurobiologically Inspired Long-Term Memory for Large Language Models. Advances in Neural Information Processing Systems 37, 2024. URL https://proceedings.neurips.cc/paper_files/paper/2024/hash/6ddc001d07ca4f319af96a3024f6dbd1-Abstract-Conference.html. Peer-reviewed graph retrieval; accessed October 7, 2026.
- [8] C. A. R. Hoare. An Axiomatic Basis for Computer Programming. Communications of the ACM, 1969. URL <https://www.cs.cmu.edu/~crary/819-f09/Hoare69.pdf>. Software foundations; accessed October 7, 2026.
- [9] Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. SWE-bench: Can Language Models Resolve Real-World GitHub Issues? International Conference on Learning Representations, 2024. URL <https://openreview.net/forum?id=VTF8yNQM66>. Peer-reviewed software benchmark; accessed October 7, 2026.
- [10] Yubin Kim, Ken Gu, Chanwoo Park, Chunjong Park, Samuel Schmidgall, A. Ali Heydari, Yao Yan, Zhihan Zhang, Yuchen Zhuang, Yun Liu, Mark Malhotra, Paul Pu Liang, Hae Won Park, Yuzhe Yang, Xuhai Xu, Yilun Du, Shwetak Patel, Tim Althoff, Daniel McDuff, and Xin Liu. Towards a Science of Scaling Agent Systems. arXiv:2512.08296v3 (2026 revision), 2025. URL <https://arxiv.org/abs/2512.08296v3>. Agent-scaling preprint; accessed October 7, 2026.
- [11] Barbara H. Liskov and Jeannette M. Wing. A Behavioral Notion of Subtyping. *ACM Transactions on Programming Languages and Systems*, 16(6):1811–1841, 1994. doi: 10.1145/197320.197383. URL <https://www.cs.cmu.edu/~wing/publications/LiskovWing94.pdf>. Software foundations; accessed October 7, 2026.
- [12] H. Penny Nii. Blackboard Systems: The Blackboard Model of Problem Solving and the Evolution of Blackboard Architectures. AI Magazine, 1986. URL <https://doi.org/10.1609/aimag.v7i2.537>. Coordination foundations; accessed October 7, 2026.
- [13] David L. Parnas. On the Criteria To Be Used in Decomposing Systems into Modules. *Communications of the ACM*, 15(12):1053–1058, 1972. doi: 10.1145/361598.361623. URL <https://doi.org/10.1145/361598.361623>. Software foundations; accessed October 7, 2026.

- [14] Marc Shapiro, Nuno Preguiça, Carlos Baquero, and Marek Zawirski. Conflict-free Replicated Data Types. Technical Report RR-7687, INRIA, 2011. URL <https://inria.hal.science/inria-00609399v2>. Distributed-state theory; accessed October 7, 2026.
- [15] Reid G. Smith. The Contract Net Protocol: High-Level Communication and Control in a Distributed Problem Solver. IEEE Transactions on Computers, 1980. URL <https://ieeexplore.ieee.org/document/1675516>. Coordination foundations; accessed October 7, 2026.
- [16] World Wide Web Consortium. SKOS Simple Knowledge Organization System Reference. W3C Recommendation, 2009. URL <https://www.w3.org/TR/2009/REC-skos-reference-20090818/>. Normative standard; accessed October 7, 2026.
- [17] World Wide Web Consortium. OWL 2 Web Ontology Language Profiles (Second Edition). W3C Recommendation, 2012. URL <https://www.w3.org/TR/2012/REC-owl2-profiles-20121211/>. Normative standard; accessed October 7, 2026.
- [18] World Wide Web Consortium. OWL 2 Web Ontology Language Direct Semantics (Second Edition). W3C Recommendation, 2012. URL <https://www.w3.org/TR/2012/REC-owl2-direct-semantics-20121211/>. Normative standard; accessed October 7, 2026.
- [19] World Wide Web Consortium. PROV-O: The PROV Ontology. W3C Recommendation, 2013. URL <https://www.w3.org/TR/2013/REC-prov-o-20130430/>. Normative standard; accessed October 7, 2026.
- [20] World Wide Web Consortium. Shapes Constraint Language (SHACL). W3C Recommendation, 2017. URL <https://www.w3.org/TR/2017/REC-shacl-20170720/>. Normative standard; accessed October 7, 2026.
- [21] Qingyun Wu, Gagan Bansal, Jieyu Zhang, Yiran Wu, Beibin Li, Erkang Zhu, Li Jiang, Xiaoyun Zhang, Shaokun Zhang, Jiale Liu, Ahmed Hassan Awadallah, Ryan W. White, Doug Burger, and Chi Wang. Auto-Gen: Enabling Next-Gen LLM Applications via Multi-Agent Conversation. arXiv:2308.08155v2, 2023. URL <https://arxiv.org/abs/2308.08155v2>. Agent framework preprint; accessed October 7, 2026.
- [22] John Yang, Carlos E. Jimenez, Alexander Wettig, Kilian Lieret, Shunyu Yao, Karthik Narasimhan, and Ofir Press. SWE-agent: Agent-Computer Interfaces Enable Automated Software Engineering. Advances in Neural Information Processing Systems 37, 2024. URL https://proceedings.neurips.cc/paper_files/paper/2024/hash/5a7c947568c1b1328ccc5230172e1e7c-Abstract-Conference.html. Peer-reviewed software agents; accessed October 7, 2026.